

AccuKnox wins the #1 AI Startup Award at Security BSides Bangalore 2026

AccuKnox took first place in **EmergeX**, the startup competition at Security BSides Bangalore 2026, judged in front of practitioners and security leaders at the Annual Cyber Security Conference in Whitefield. **It is the second year in a row** that AccuKnox has topped the category. What earned it was live enforcement, a booth running real prompt injection and agentic AI attacks against a working instance while attendees watched.



#1

EmergeX startup award,
BSides Bangalore ACSC
2026

2 yrs

Back to back wins, 2025
and 2026

2

Speaker sessions on the
main agenda

8


AI security modules
integrated into a single
policy engine

"The depth of the AccuKnox Agentic AI Security platform stood out even more than the first time around. Back to back wins at Security BSides Bangalore do not happen by chance. They happen when a team keeps building instead of pausing to celebrate the last trophy."

Sujatha Yakasiri, Founder, Security BSides Bangalore

A working control plane, running live at the booth

Attendees queued up to attack the platform rather than to watch a slide deck. Eight AI security modules run on the same policy engine AccuKnox uses for cloud workloads, all reachable from one console.



Eight AI security modules over AI assets, substrate, and every deployment model, including air-gapped and edge.

Enforcement lives inside the kernel

KubeArmor, the open source project AccuKnox maintains, uses **eBPF and LSM** to allow or block process, file, and network behavior inside the kernel. The same engine that guards production workloads now guards AI agents and MCP servers.

Prompt defense that holds state across turns

Classifier-only prompt firewalls have a measurable bypass problem. Published research records **100% attack success** for some character injection techniques, and multi-turn frameworks such as HarmNet and X-Teaming land at **94 to 99%**. AccuKnox tracks context across turns, inspects input and output, masks PII and PHI, and catches secrets before they leave the session.



Agents get an identity and a permission boundary

Identity comes from **SPIFFE**, authorization from **OpenFGA**, isolation from the kernel. An agent that reaches for a tool it never needed gets stopped rather than logged after the fact.

Red teaming runs on every model update

Automated adversarial probes for injection, jailbreak, and hallucination fire after each change, because a model that ships weekly cannot be signed off once a year. Findings map to **MITRE ATLAS, OWASP LLM Top 10, NIST AI RMF, and the EU AI Act**.

Deployment parity is the part regulated buyers test first. One policy engine, one audit trail, and identical scanning across SaaS, Kubernetes, private cloud, and fully air-gapped installs. Coverage reaches cloud unmanaged and on-prem unmanaged AI too, which is where shadow AI hides. IBM priced that blind spot at **\$670K** in extra average breach cost in 2025.



Rahul Jadhav

CTO AND CO-FOUNDER

The Subtle Art of AI Prompt Guardrails

Where guardrails hold, where they break, and what the literature actually shows. Attack techniques with their measured success rates, the latency cost of inspection, and the integration patterns that survive production.

Gaurav Mishra

SENIOR PRODUCT MANAGER

A multi-layered defence pipeline for an enterprise LLM chat application

One real enterprise chat application, traced through the full defence chain, from prompt inspection to dataset controls to runtime enforcement on the serving infrastructure.

Why AccuKnox took first place

The judges could attack it themselves

Prompt injection and agentic attacks ran against a live instance, and each block landed in the audit trail while the attacker watched.

The enforcement claim is auditable

KubeArmor is open source with a public codebase and an active contributor base, so the kernel-level enforcement story can be read rather than taken on trust.

Coverage includes the assets nobody wants to discuss

Rogue notebooks, self-hosted models on VMs, developer workstation LLMs, and MCP servers sit under one policy engine with the sanctioned cloud services.

What this means for you

For organizations across financial services, healthcare, telecommunications, manufacturing, critical infrastructure, government, and other regulated industries, AI security is now a standard part of every customer and compliance review. AccuKnox addresses these requirements with a unified platform that deploys wherever sensitive data must remain, enforces protection at runtime to stop threats before they become incidents, and generates the evidence auditors and customers expect. Recognition at BSides validates the platform's capabilities through independent testing by security practitioners in real-world environments.