



Azure Cloud Security and AI Posture Report

30

Azure subscriptions

11,221

Cloud assets

31,439

Active findings

Certified & Accredited by



As Featured In



Available On



PART 1


Cloud Security Posture Management



30 Azure subscriptions, 11,221 assets and 31,439 active findings, read through exposure and attack paths.

POSTURE VOLUME, EXPOSURE AND COMPOUND RISK

A short view of where the confirmed Azure footprint carries risk today.



FINDING SEVERITY DISTRIBUTION



● High 10,144 ● Medium 13,353 ● Low 7,942

74.7%

of active findings are Medium or High severity

The estate needs broad risk reduction. Drive remediation by exploitability and attack paths rather than raw finding volume.

PRIORITY RISK SIGNALS

- **External attack surface**
412 confirmed Internet-exposed workloads, plus Internet-open management protocols.
- **Data protection concentration**
Storage and block storage dominate High finding volume. Several controls affect public access and encryption.
- **Identity and trust blast radius**
Confirmed toxic paths involve domain-controller and Root CA systems exposed through permissive NSGs.

WHAT THE ASSESSMENT COVERED

The estate is dominated by network, compute and block-storage assets, the primary security control surface.

IMPORTANT ASSET DISTRIBUTION



11,221

TOTAL CONFIRMED AZURE ASSETS

● Network plane

3,385 networking assets make connectivity, segmentation and logging a major attack-surface concern.

● Workload plane

Compute and host assets represent 3,035 workloads where exposure and hardening drive exploitability.

● Data plane


2,483 block and object storage assets create concentrated public-access and encryption risk.

Coverage note: IAM and Cloud Account categories returned 0 assets in the queried distribution, despite account-level CSPM findings.

WHERE THE HIGH FINDINGS SIT

High findings concentrate in storage and compute controls, which gives a narrow remediation surface with broad impact.

SEVERITY VOLUME



0 confirmed Critical findings

Compound-risk override

Severity alone understates risk. Confirmed toxic combinations raise several High controls to Critical remediation priority.

HIGH FINDINGS BY ASSET CATEGORY



90.5%

of all High findings sit in Block Storage, Compute and Object Storage.

Technical interpretation


Fix the systemic controls first: private access by default, encryption at host, BYOK and CMK, and exposure guardrails. Do policy cleanup after.

Highest-Volume High Findings

THE CONTROLS BEHIND THE VOLUME

The top controls point to data-protection and exposure weaknesses, followed by compute hardening gaps.

TOP HIGH FINDINGS



TECHNICAL INTERPRETATION

Storage exposure

Disk and snapshot controls are the largest single source of High findings. Fix them as a platform baseline.

Encryption posture

BYOK and CMK controls show real gaps across block and object storage, which affects data governance.

Internet reachability

Internet Exposure is lower volume but far higher exploitability. It must override volume-based ranking.

AI data controls

OpenAI and ML workspaces also show CMK and logging gaps, which creates governance risk for AI workloads.

INTERNET REACHABILITY AND OPEN PORTS

Permissive NSG rules and publicly reachable administrative services amplify Internet reachability.



PUBLIC AND INTERNET EXPOSURE SIGNALS



CONTROL-PLANE GAPS THAT RAISE EXPLOITABILITY

- Network observability**
956 NSG Flow Log gaps and 954 NSG Log Analytics gaps cut visibility into attacker movement and egress.
- VNET and load-balancer logging**
VNET logging gaps and 124 load-balancer Log Analytics gaps reduce forensic depth for exposed services.
- Key management exposure**
13 public Key Vault findings, plus missing private endpoints and restrictive defaults, widen secret-store reach.

High-Risk Publicly Exposed Assets

CONFIRMED ASSET TO NETWORK RELATIONSHIPS

These systems are where exposure carries immediate attack-path meaning.

ASSET	TYPE	CONFIRMED EXPOSURE AND RELATIONSHIP	RISK
ZPW-IGA-DC01	Azure VM	Internet-exposed via NSG ZPWIGADC01nsg882. The related NSG has all ports open to *	CRITICAL
ZPW-IGA-DC02	Azure VM	Internet-exposed via NSG ZPW-IGA-DC02-nsg. All ports open to *	CRITICAL
ZPW-DEV-DC01	Azure VM	Internet-exposed via ZPW-DEV-DC01-NSG. All ports open to *	CRITICAL
IGROOTCA	Azure VM	Internet-exposed via IGROOTCA-NSG. SMB TCP 445 open to *	CRITICAL
AzureCloud-OpenShiftBastionHost	Azure VM	Internet-exposed. RDP TCP 3389 open to *	HIGH/CRIT
ocpintprdcldu2-kgz68-master-2	Azure VM	Internet-exposed through public load balancer ocpintprdcldu2-kgz68	HIGH
ZPW-MFT-PGWS01	Azure VM	Internet-exposed. RDP and SSH open to *	HIGH
ZPW-MFT-PGWS02	Azure VM	Internet-exposed. RDP and SSH open to *	HIGH
medical-records	Storage	Public blob access and no CMK encryption	CRITICAL
KEYVAULT-CIAM-PRD	Key Vault	Public access, no private endpoint, default network allows all networks	CRITICAL

Toxic Combination Matrix

FIVE CONFIRMED MULTI-CONTROL PATHS


Each row is a place where separate posture failures combine into materially higher risk.

ID	TOXIC COMBINATION	CONFIRMED RELATIONSHIP	POTENTIAL BLAST RADIUS	RISK
TC-01	Internet-exposed DCs plus all ports open	VM exposure names the same NSG, with all protocols and all ports open to *	Identity compromise, lateral movement	CRITICAL
TC-02	Internet-exposed Root CA plus SMB open	IGROOTCA is exposed through IGROOTCA-NSG. TCP 445 is open to *	Trust and certificate infrastructure compromise	CRITICAL
TC-03	Public Key Vault plus no private endpoint	The same Key Vault has public access, no private endpoint, an all-network default and a logging gap	Secrets, keys and certificates	CRITICAL
TC-04	Public medical-records container plus no CMK	The same storage container is public and is not CMK-encrypted	Sensitive data exposure and governance	CRITICAL
TC-05	Internet-facing OpenShift bastion plus RDP	VM exposure names the NSG. RDP 3389 is open to *. Hardening and logging gaps are present	Pivot into OpenShift and internal workloads	HIGH/CRIT

Identity and Trust Plane Attack Paths

TC-01 AND TC-02 IN DETAIL

Confirmed exposure chains show how Internet reachability meets high-value identity and trust infrastructure.



● Security implication

The network path is explicitly confirmed to the named high-value systems, which raises blast radius. Treat these as connected exposures rather than isolated misconfigurations.

TC-03 AND TC-04 IN DETAIL

Public service endpoints combine with weak network isolation, encryption or logging controls.

TC-03

KEYVAULT-CIAM-PRD

CRITICAL

KEY VAULT

Public Key Vault endpoint
Secrets, keys and certificates

No private endpoint
No private-only service path

Default access is all networks
Broad network reachability

Logging gap
Key Vault Log Analytics is not enabled

Potential impact: Reduced detection, plus secret exposure risk

TC-04

medical-records

CRITICAL

STORAGE CONTAINER

Public blob container
Reachable object data named medical-records

No CMK control
Customer-managed key is absent

Sensitive-data context
Asset naming implies high sensitivity

Public exposure
Internet-accessible object data

Potential impact: Data exposure, plus weaker key governance

Risk Concentration by Subscription and Region

WHERE TO SEND THE EFFORT

A small subset of subscriptions, and Central India, account for most of the observed CSPM risk.

TOP SUBSCRIPTIONS BY HIGH FINDINGS



Top account: 2,782 High findings and 7,065 total active findings.

FINDINGS BY REGION




Regional hotspot and ranking rule

Central India carries 27,526 total findings and 9,105 High findings, which makes it the primary remediation domain. Pair account volume with toxic-combination context. A lower-volume account can still outrank a larger one when it holds exposed crown-jewel assets.

FAILED MAPPINGS AND THE PLAN

Failed framework mappings show the control themes. Remediation is ranked by attack-path impact rather than mapping volume.

TOP FAILED COMPLIANCE MAPPINGS



Note: compliance data represents failed finding and control mappings. It is not a complete pass percentage.

PRIORITIZED REMEDIATION PLAN

- P0 Break confirmed attack paths**
Restrict NSGs to approved sources and ports. Remove public access from the DC, Root CA, Key Vault and storage paths.
- P1 Reduce external exposure**
Remove unnecessary public IP and service exposure. Close RDP, SSH and SMB. Enforce private endpoints and restrictive defaults.
- P1 Close data-protection gaps**
Enforce CMK and BYOK where required, disk encryption at host, and private storage access.
- P2 Restore observability**
Enable NSG, VNET, load-balancer and Key Vault diagnostics, with central Log Analytics integration.
- P2 Baseline identity controls**
Enable managed identities where applicable. Standardize identity-centric workload hardening.

EVIDENCE VERSUS GAPS

The report separates confirmed evidence from context that was not available or not validated.

EVIDENCE USED

Cloud provider

Azure

Source

Cloud Findings

Filters

Active, not ignored, azure_subscription

Confirmed Azure accounts

30 with active findings

Toxic-combination rule

Same asset or finding explicitly naming the related NSG or load balancer

NOT CONFIDENTLY AVAILABLE OR VALIDATED

Public IP addresses

Not returned in the queried CSPM fields

Excessive IAM and privilege escalation

No confirmed Azure findings retrieved

Secret-scan and container-secret findings

No active Azure findings in the queried data

Total onboarded Azure subscriptions

36 filter values exist. Only 30 carry confirmed active CSPM findings

Compliance score

Mappings available. No complete pass or fail percentage returned

● Reporting principle

Only confirmed evidence supports the attack-path assertions. Missing context is surfaced rather than inferred.

PART 2 · AI AND ML SECURITY POSTURE

Azure AI Security Posture Report

AI asset inventory, exposure, security findings and reporting confidence across the confirmed Azure AI and ML estate.

CONFIRMED SCOPE

34

AI and ML assets
Inventory-confirmed

164

Active AI findings
Across 34 assets

24

Public network access
Confirmed endpoints

Assessment basis: confirmed AI asset inventory, active findings, and no unvalidated account-level asset counts

AI POSTURE SIGNAL SUMMARY

24

Exposure

Assets with public network access

51

Severity

High-severity AI findings

41

Data

Data-security findings

2

Platform


Azure ML and Databricks

Executive AI Risk Snapshot



AI VOLUME, SEVERITY AND EXPOSURE

A short view of the confirmed Azure AI and ML footprint.



FINDING SEVERITY DISTRIBUTION



● High 51

● Medium 96

● Low 17

89.6%

of AI findings are Medium or High severity

Risk concentrates in Azure Machine Learning controls: public access, customer-managed encryption, diagnostics and monitoring.

PRIORITY RISK SIGNALS

● Public network exposure

24 AI assets have public network access enabled. No separate Internet Exposure finding was returned.

● Data protection concentration

All 34 assets carry data-security findings. 41 data-security findings were confirmed across the AI estate.

● Platform concentration

27 of 34 assets are Azure Machine Learning workspaces, and all 27 carry High-severity findings.

PLATFORM AND REGION SPREAD

The confirmed AI estate is entirely Azure. It concentrates in Machine Learning workspaces deployed mainly in southindia.

AI SERVICE AND PLATFORM DISTRIBUTION



REGIONAL ASSET DISTRIBUTION



Coverage note: unique AI asset counts by account, subscription or project were absent from the inventory output. They are not inferred here.

34

TOTAL CONFIRMED AI AND ML ASSETS

Platform concentration

79.4% of AI assets are Azure Machine Learning workspaces.

Regional concentration

64.7% of AI assets sit in southindia.

Finding coverage

34 of 34 confirmed AI assets carry at least one active finding.

WHICH CONTROLS CARRY THE WEIGHT

High-severity impact is broad across Azure ML workspaces. Medium and Low findings touch the full AI estate.

SEVERITY VOLUME




0 confirmed Critical findings

Coverage interpretation

High severity touches 27 unique assets.
Medium and Low each touch all 34.

TOP AI FINDINGS



89.6%

of all AI findings are Medium or High severity.

PUBLIC ACCESS AND CONTROL PATTERNS

Public network access is the dominant exposure signal. Data-security controls affect the whole AI estate.

24

PUBLIC NETWORK ACCESS

24 confirmed findings

0

INTERNET EXPOSURE FINDINGS

None separately confirmed

0

IAM AND IDENTITY FINDINGS

None confirmed

41

DATA-SECURITY FINDINGS

Across 34 unique assets

EXPOSURE INTERPRETATION

Public endpoint enabled

24 unique AI assets have public network access enabled, based on confirmed ML Workspace findings.

Public access is not Internet Exposure

No AI-specific finding named Internet Exposure was returned. Public network reachability is reported on its own.

Identity scope boundary

No confirmed AI IAM or identity posture findings came back from the available dataset.

DATA SECURITY CONTROL PATTERNS

ML Workspace CMK encryption 22

Databricks Managed Services CMK 7

Databricks Managed Disk CMK 7

Diagnostic logging and monitoring 21

34

unique AI assets carry data-security findings

Azure ML AGAINST DATABRICKS

Azure Machine Learning is most of the AI estate and carries the whole confirmed High-affected asset population.

PLATFORM POSTURE COMPARISON

Azure Machine Learning Workspace

27

Unique assets

141

Total findings

27

High-affected

Azure Databricks Workspace

7

Unique assets

23


Total findings

0

High-affected

Estate split: Azure ML is 79.4% of assets, Databricks is 20.6%.

SECURITY-DOMAIN COVERAGE



● Priority

Reduce public network access first, then strengthen encryption and logging controls on the Azure ML workspaces.

WHERE THE AI FINDINGS CLUSTER


AI finding volume sits mostly in one account. Asset inventory clusters in southindia.

AI FINDING-COUNT DISTRIBUTION BY ACCOUNT



The top account contributes 129 of 164 active AI findings, or 78.7%.

UNIQUE AI ASSETS BY REGION



Regional concentration

22 of 34 assets, or 64.7%, sit in southindia.

Account inventory boundary

Unique AI asset counts by account, subscription or project could not be validated from the inventory output. The account view here is finding-count distribution only.

WHAT THE COUNTS DO AND DO NOT SAY

Every count below keeps the source-data boundary. No account-level asset relationship is inferred.

VALIDATION NOTES

Inventory basis

34 unique AI asset IDs in the confirmed inventory output

Account limitation

Account, subscription and project were absent from the inventory output

Account view

The account distribution here is active AI finding-count distribution only

Public endpoint basis

24 public-network-access assets come from confirmed ML Workspace findings

Internet exposure

No separate AI-specific Internet Exposure findings were returned

Identity scope

No confirmed AI IAM or identity findings were returned

KEY COUNT-BASED HIGHLIGHTS

34 Total AI and ML assets discovered

27 Azure Machine Learning workspaces

7 Azure Databricks workspaces

27 AI assets with High findings

24 AI assets with public network access

0 Critical AI findings confirmed

0 AI IAM and identity findings confirmed

● Reporting principle

Only confirmed tenant data appears here. Unsupported values stay unreported.

SAME-ASSET RISK COMBINATIONS

Five high-confidence combinations are confirmed. No complete multi-hop path to identity, data, models or storage was proven.

5

TOXIC COMBINATIONS

All include High controls

4

START AT PUBLIC ACCESS

Azure ML workspace exposure

5

HIGH-RISK COMBINATIONS

High findings involved

0

MULTI-HOP PATHS

No downstream target proven

0

PRIVILEGED IDENTITY PATHS

No privileged role confirmed

PRIMARY AI ASSET	CONFIRMED WEAKNESS STACK	RISK	HOPS
whatsapp-chat	Public access, no CMK, no diagnostics	HIGH	1
6eskai-innov8	Public access, no CMK, HBI disabled, Basic SKU	HIGH	1
6e-enterprise	Public access, no CMK, no diagnostics	HIGH	1
AI-Foundry-NP-EA	Public access, no CMK, HBI disabled	HIGH	1
6e-openai-Non-prod-southindia-3	No CMK, no managed identity, no diagnostics	MED/HIGH	1

Confirmed

Public or API entry, plus same-asset configuration findings.


Not confirmed

Privileged identities, datasets, model endpoints, credentials, storage or lateral movement.

Source: confirmed AI toxic-combination analysis, same-asset finding relationships

FOUR CONFIRMED COMBINATIONS

Four paths start with public network access, then stack encryption, logging and governance weaknesses on the same workspace.



Evidence boundary

All four paths rest on the public-workspace finding and same-asset control findings. No identity, dataset, model endpoint, storage or credential relationship was confirmed.

THE FIFTH COMBINATION

Combination five affects an Azure OpenAI and Cognitive account. The aggregate view shows recurring encryption, identity and logging gaps.

TOXIC COMBINATION 5, AZURE OPENAI AND COGNITIVE

6e-openai-Non-prod-southindia-3

MEDIUM / HIGH

azure_cognitive_account

HIGH

OpenAI Account CMK Encrypted

The account is not encrypted with a customer-managed key

MEDIUM

OpenAI Account Managed Identity Enabled

Managed identity is not enabled

MEDIUM

OpenAI Account Diagnostic Logging Enabled

Diagnostic logs are not enabled

Public exposure, privileged identity, model deployment, dataset and credential relationships were not confirmed for this asset.

AGGREGATE CONFIRMED PATTERNS

	Public Azure ML workspace plus CMK gap	24 findings
	Azure ML workspace plus diagnostics disabled	19 findings
	Azure ML workspace plus HBI disabled	27 findings
	Azure OpenAI or Cognitive plus CMK disabled	82 findings
	Azure OpenAI or Cognitive plus no managed identity	40 findings
	Azure OpenAI or Cognitive plus diagnostics disabled	75 findings

WHAT WAS NOT CONFIRMED

0

complete multi-hop paths

0

privileged identity paths

0

dataset and data-source paths

0

model-endpoint exposure paths

0

credential exposure paths



SEE US IN ACTION

support@accuknox.com

Certified by



www.AccuKnox.com

